

In [1]: `!pip install joblib`

Requirement already satisfied: joblib in c:\users\thinithi\anaconda3\lib\site-packages (1.2.0)

```
In [1]: import os # cheeck currently directory
import pandas as pd # Load pandas library for data wrangling
# to convert weather station to a numeric number
from sklearn.preprocessing import LabelEncoder
# Split dataset and conduct cross-validation
from sklearn.model_selection import train_test_split, cross_val_score
# build a random forest model
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score
from IPython.display import display
from sklearn.preprocessing import StandardScaler # Standardize column
# build a logisitc regression model
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier # build a decision tree model
# build a gradient boosting model
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier # build a knn model
from sklearn.svm import SVC # Build a svm model
from sklearn.naive_bayes import GaussianNB # build naive bayes model
from collections import Counter # obtain counts
import matplotlib.pyplot as plt # plot correlation matrix
import seaborn as sns # plot feature importance
import joblib # save model
# utilize calss weight balancing in models
from sklearn.utils.class_weight import compute_class_weight
from PIL import Image # Import model evaluation report image
from sklearn.model_selection import GridSearchCV # conduct hyper parameter tuning
```

```
In [ ]: # Get the current working directory
current_directory = os.getcwd()
current_directory
```

Data Exploration

```
In [2]: file_path = 'Bushfire_Weather_data.xlsx'

# Read the Excel file into a DataFrame
df = pd.read_excel(file_path)
df.head()
```

Out[2]:

	Page_URL	Page_title	Date	Day	Temps_Min	Temps_Max	Rain	AM9_Temp	AM9_RH	AM9_D
0	201401	Benalla	1	We	9.8	26.9	0.0	21.1	37	S
1	201401	Benalla	2	Th	14.9	25.9	3.6	19.0	72	SSV
2	201401	Benalla	3	Fr	15.8	30.2	0.0	22.0	35	SS
3	201401	Benalla	4	Sa	11.4	25.0	0.0	15.5	51	SSV
4	201401	Benalla	7	Tu	9.2	25.9	0.0	15.0	43	

In [3]: `df.count()`

```
Out[3]: Page_URL      20021
Page_title  20021
Date        20021
Day         20021
Temps_Min   20021
Temps_Max   20021
Rain        20021
AM9_Temp    20021
AM9_RH      20021
AM9_Dir     20015
AM9_Spd     20021
PM3_Temp    20021
PM3_RH      20021
PM3_Dir     20011
PM3_Spd     20021
KEY         20021
dtype: int64
```

```
In [4]: file_path = 'Bushfire_Wstation_2016_2023_01.csv'

# Read the CSV file into a DataFrame with a different encoding
df2 = pd.read_csv(file_path, encoding='ISO-8859-1')

df2.head()
```

```
Out[4]:
```

	FIRETYPE	SEASON	FIRE_NO	NAME	START_DATE	STRTDATIT	TREAT_TYPE	FIREKE
				BARANDUDA				
				-				
0	BUSHFIRE	2017	32	BARANDUDA RANGE TRACK	3/2/2017	20170302	FIRE	W20173403
1	BUSHFIRE	2017	32	GLENROWAN RIDGE TRACK	12/20/2016	20161220	FIRE	W20173703
2	BUSHFIRE	2016	61	MT BRUNO - JONES RD	12/20/2015	20151220	FIRE	W20163706
3	BUSHFIRE	2018	49	OVENS 49 - BEECHWORTH - WARNERS TRACK	1/23/2018	20180123	FIRE	W20183404
4	BUSHFIRE	2016	2	BAILIESTON - REEDY LAKE	8/18/2015	20150818	FIRE	W201699

```
In [5]: df2.count()
```

```
Out[5]: FIRETYPE      519
        SEASON        519
        FIRE_NO       516
        NAME          519
        START_DATE    519
        STRTDATIT     519
        TREAT_TYPE    519
        FIREKEY       493
        DSE_ID        507
        CFA_ID        81
        AREA_HA       519
        latitude      519
        longitude     519
        WStation      519
        Year          519
        Month         519
        Day           519
        dtype: int64
```

```
In [6]: # Rename Wstation
df['Page_title'].replace('Wangaratta', 'WANGARATTA', inplace=True)
df['Page_title'].replace('Benalla', 'BENALLA', inplace=True)
df['Page_title'].replace('Shepparton', 'SHEPPARTON', inplace=True)
df['Page_title'].replace('Falls Creek', 'FALLS_CREEK', inplace=True)
df['Page_title'].replace('Mount Buller', 'MOUNT_BULLER', inplace=True)
df['Page_title'].replace('Yarrawonga', 'YARRAWONGA', inplace=True)
df['Page_title'].replace('Albury', 'ALBURY', inplace=True)
df['Page_title'].replace('Corryong', 'CARRYONG', inplace=True)
df['Page_title'].replace('Edi Upper', 'EDI_UPPER', inplace=True)
df['Page_title'].replace('Hunters Hill', 'HUNTERS_HILL', inplace=True)
df['Page_title'].replace('Rutherglen', 'RUTHERGLEN', inplace=True)
```

```
In [7]: # Rename Wstation
df2['WStation'].replace('WANGARATTA AERO', 'WANGARATTA', inplace=True)
df2['WStation'].replace('BENALLA AIRPORT', 'BENALLA', inplace=True)
df2['WStation'].replace('SHEPPARTON AIRPORT', 'SHEPPARTON', inplace=True)
df2['WStation'].replace('FALLS CREEK', 'FALLS_CREEK', inplace=True)
df2['WStation'].replace('MOUNT BULLER', 'MOUNT_BULLER', inplace=True)
df2['WStation'].replace('YARRAWONGA', 'YARRAWONGA', inplace=True)
df2['WStation'].replace('ALBURY AIRPORT AWS', 'ALBURY', inplace=True)
df2['WStation'].replace('CARRYONG AIRPORT', 'CARRYONG', inplace=True)
df2['WStation'].replace('EDI UPPER', 'EDI_UPPER', inplace=True)
df2['WStation'].replace('HUNTERS HILL', 'HUNTERS_HILL', inplace=True)
df2['WStation'].replace('RUTHERGLEN RESEARCH', 'RUTHERGLEN', inplace=True)
```

```
In [8]: df = df.drop(['Page_URL', 'Date', 'Day'], axis=1)
df.rename(columns = {'Page_title': 'Station_Region'}, inplace = True)
df.head()
```

```
Out[8]:
```

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_Temp	AM9_RH	AM9_Dir	AM9_Spd	PM3
0	BENALLA	9.8	26.9	0.0	21.1	37	SE	7.0	
1	BENALLA	14.9	25.9	3.6	19.0	72	SSW	4.0	
2	BENALLA	15.8	30.2	0.0	22.0	35	SSE	4.0	
3	BENALLA	11.4	25.0	0.0	15.5	51	SSW	6.0	
4	BENALLA	9.2	25.9	0.0	15.0	43	S	19.0	

```
In [9]: # Summarizing the DataFrame
summary_df = df2.groupby(['WStation', 'STRTDATIT']).agg(
    total_area_ha=('AREA_HA', 'sum'),
    count_of_rows=('AREA_HA', 'size')
).reset_index().sort_values(by='count_of_rows', ascending=False)

# Display the summary DataFrame
summary_df.head()
```

```
Out[9]:
```

	WStation	STRTDATIT	total_area_ha	count_of_rows
168	MOUNT_BULLER	20191231	83.4484	41
57	DARTMOUTH RESERVOIR	20191121	6282.5712	31
118	FALLS_CREEK	20191231	35830.2541	21
119	FALLS_CREEK	20200101	44315.9390	14
110	FALLS_CREEK	20190228	13236.3868	14

```
In [10]: df['Station_Region'].unique()
```

```
Out[10]: array(['BENALLA', 'WANGARATTA', 'YARRAWONGA', 'SHEPPARTON', 'FALLS_CREEK',
        'MOUNT_BULLER', 'RUTHERGLEN', 'ALBURY', 'CARRYONG', 'EDI_UPPER',
        'HUNTERS_HILL'], dtype=object)
```

```
In [11]: summary_df['WStation'].unique()
```

```
Out[11]: array(['MOUNT_BULLER', 'DARTMOUTH RESERVOIR', 'FALLS_CREEK', 'EDI_UPPER',
        'CARRYONG', 'ALBURY', 'HUNTERS_HILL', 'STRATHBOGIE', 'SHEPPARTON',
        'YARRAWONGA', 'BENALLA', 'RUTHERGLEN', 'WANGARATTA'], dtype=object)
```

```
In [12]: # Convert the `KEY` column in df to string
df['KEY'] = df['KEY'].astype(str)

# Create a new column by concatenating WStation and STRTDATIT
summary_df['JOINKEY1'] = summary_df['WStation'] + '_' + summary_df['STRTDATIT'].ast
df['JOINKEY2'] = df['Station_Region'] + '_' + df['KEY']

# Left join df (weather data) with df2 (fire data) on KEY and STRTDATIT
data = df.merge(summary_df, how='left', left_on='JOINKEY2', right_on='JOINKEY1')
data.head()
```

```
Out[12]:
```

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_Temp	AM9_RH	AM9_Dir	AM9_Spd	PM3
0	BENALLA	9.8	26.9	0.0	21.1	37	SE	7.0	
1	BENALLA	14.9	25.9	3.6	19.0	72	SSW	4.0	
2	BENALLA	15.8	30.2	0.0	22.0	35	SSE	4.0	
3	BENALLA	11.4	25.0	0.0	15.5	51	SSW	6.0	
4	BENALLA	9.2	25.9	0.0	15.0	43	S	19.0	

```
In [13]: # Create a new column based on the value of count_of_rows
data['Bushfire_Event_Target'] = data['count_of_rows'].apply(lambda x: 0 if pd.isna(
data.head()
```

Out[13]:

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_Temp	AM9_RH	AM9_Dir	AM9_Spd	PM3
0	BENALLA	9.8	26.9	0.0	21.1	37	SE	7.0	
1	BENALLA	14.9	25.9	3.6	19.0	72	SSW	4.0	
2	BENALLA	15.8	30.2	0.0	22.0	35	SSE	4.0	
3	BENALLA	11.4	25.0	0.0	15.5	51	SSW	6.0	
4	BENALLA	9.2	25.9	0.0	15.0	43	S	19.0	

In [14]:

```
# Remove unwanted columns from the DataFrame
data = data.drop(columns=['WStation'])
data = data.drop(columns=['STRDATIT'])
data = data.drop(columns=['count_of_rows'])
data = data.drop(columns=['AM9_Dir'])
data = data.drop(columns=['PM3_Dir'])
data = data.drop(columns=['KEY'])
data = data.drop(columns=['JOINKEY1'])
data = data.drop(columns=['JOINKEY2'])

# Replace NaN values in the 'Total_Area' column with 0
data['total_area_ha'] = data['total_area_ha'].fillna(0)
data.rename(columns = {'total_area_ha': 'Total_Area'}, inplace = True)
# Normalize the Total_Area column
data['Total_Area'] = (data['Total_Area'] - data['Total_Area'].min()) / (data['Total_Area'].max() - data['Total_Area'].min())
```

In [15]: data.head()

Out[15]:

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_Temp	AM9_RH	AM9_Spd	PM3_Temp	PI
0	BENALLA	9.8	26.9	0.0	21.1	37	7.0	25.8	
1	BENALLA	14.9	25.9	3.6	19.0	72	4.0	24.4	
2	BENALLA	15.8	30.2	0.0	22.0	35	4.0	25.5	
3	BENALLA	11.4	25.0	0.0	15.5	51	6.0	22.9	
4	BENALLA	9.2	25.9	0.0	15.0	43	19.0	23.9	

In [16]:

```
# Initialize the encoder
label_encoder = LabelEncoder()

# Fit and transform the data
data["Station_Region"] = label_encoder.fit_transform(data["Station_Region"])

data.head()
```

Out[16]:

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_Temp	AM9_RH	AM9_Spd	PM3_Temp	PM3_RH	PM3_Spd	Total_Area	Bushfire_Event_Target
0	1	9.8	26.9	0.0	21.1	37	7.0	25.8	51	19.0	23.9	0.0
1	1	14.9	25.9	3.6	19.0	72	4.0	24.4	51	19.0	23.9	0.0
2	1	15.8	30.2	0.0	22.0	35	4.0	25.5	51	19.0	23.9	0.0
3	1	11.4	25.0	0.0	15.5	51	6.0	22.9	51	19.0	23.9	0.0
4	1	9.2	25.9	0.0	15.0	43	19.0	23.9	51	19.0	23.9	0.0

In [17]: *# Print the mapping of original values to encoded values*

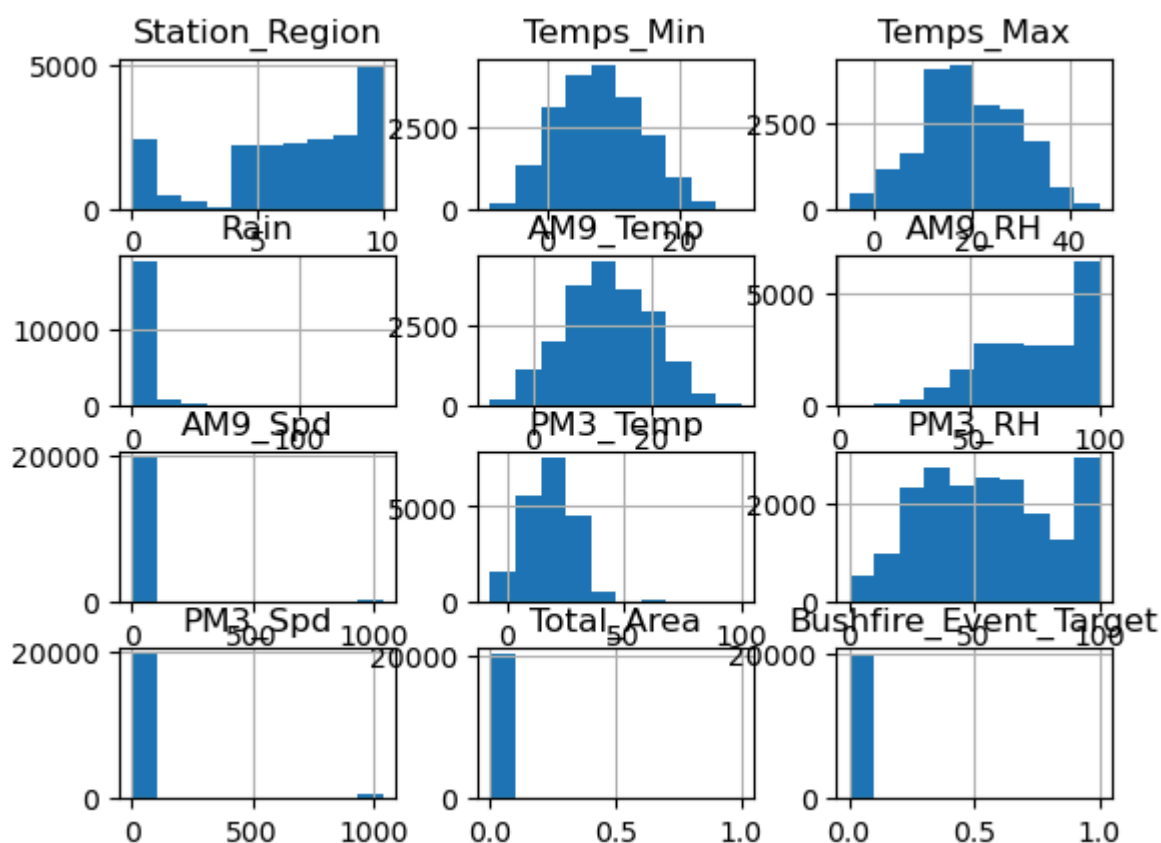
```
mapping = dict(zip(label_encoder.classes_, label_encoder.transform(label_encoder.classes_)))
print("Mapping", mapping)
```

Mapping {'ALBURY': 0, 'BENALLA': 1, 'CARRYONG': 2, 'EDI_UPPER': 3, 'FALLS_CREEK': 4, 'HUNTERS_HILL': 5, 'MOUNT_BULLER': 6, 'RUTHERGLEN': 7, 'SHEPPARTON': 8, 'WANGARATTA': 9, 'YARRAWONGA': 10}

In [18]: *# Assigning the data to the DataFrame*

```
df = pd.DataFrame(data)
df.hist()
```

Out[18]: array([[<Axes: title={ 'center': 'Station_Region' }>,
<Axes: title={ 'center': 'Temps_Min' }>,
<Axes: title={ 'center': 'Temps_Max' }>],
[<Axes: title={ 'center': 'Rain' }>,
<Axes: title={ 'center': 'AM9_Temp' }>,
<Axes: title={ 'center': 'AM9_RH' }>],
[<Axes: title={ 'center': 'AM9_Spd' }>,
<Axes: title={ 'center': 'PM3_Temp' }>,
<Axes: title={ 'center': 'PM3_RH' }>],
[<Axes: title={ 'center': 'PM3_Spd' }>,
<Axes: title={ 'center': 'Total_Area' }>,
<Axes: title={ 'center': 'Bushfire_Event_Target' }>]], dtype=object)



```
In [ ]: df.head()
```

```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20021 entries, 0 to 20020
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Station_Region        20021 non-null  int32
 1   Temps_Min              20021 non-null  float64
 2   Temps_Max              20021 non-null  float64
 3   Rain                   20021 non-null  float64
 4   AM9_Temp               20021 non-null  float64
 5   AM9_RH                 20021 non-null  int64
 6   AM9_Spd                20021 non-null  float64
 7   PM3_Temp               20021 non-null  float64
 8   PM3_RH                 20021 non-null  int64
 9   PM3_Spd                20021 non-null  float64
10   Total_Area              20021 non-null  float64
11   Bushfire_Event_Target  20021 non-null  int64
dtypes: float64(8), int32(1), int64(3)
memory usage: 1.8 MB
```

Resolving Class Imbalance Issues

```
In [20]: # Separate the majority and minority classes
X_0 = df[df['Bushfire_Event_Target'] == 0] # Majority class
X_1 = df[df['Bushfire_Event_Target'] == 1] # Minority class

# Randomly sample from the majority class
n_samples = 200 # Number of samples to keep from the majority class
X_0_sampled = X_0.sample(n=n_samples, random_state=42)

# Combine the undersampled majority class with the minority class
df = pd.concat([X_0_sampled, X_1])
```

```
In [21]: a = df[df['Bushfire_Event_Target'] == 0]
a.count()
```

```
Out[21]: Station_Region        200
Temps_Min              200
Temps_Max              200
Rain                   200
AM9_Temp               200
AM9_RH                 200
AM9_Spd                200
PM3_Temp               200
PM3_RH                 200
PM3_Spd                200
Total_Area              200
Bushfire_Event_Target  200
dtype: int64
```

```
In [22]: a = df[df['Bushfire_Event_Target'] == 1]
a.count()
```

```
Out[22]: Station_Region      140  
         Temps_Min          140  
         Temps_Max          140  
         Rain               140  
         AM9_Temp           140  
         AM9_RH             140  
         AM9_Spd            140  
         PM3_Temp           140  
         PM3_RH             140  
         PM3_Spd            140  
         Total_Area         140  
         Bushfire_Event_Target 140  
         dtype: int64
```


Check for correlation among the variables

```
In [23]: # Calculate the correlation matrix  
         corr_matrix = df.corr()  
  
         # Print the correlation matrix  
         print(corr_matrix)  
  
         # Visualize the correlation matrix using a heatmap  
         plt.figure(figsize=(8, 6))  
         sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)  
         plt.title('Correlation Matrix')  
         plt.show()
```


	Station_Region	Temps_Min	Temps_Max	Rain	\
Station_Region	1.000000	0.056108	0.102297	-0.044611	
Temps_Min	0.056108	1.000000	0.817511	-0.049770	
Temps_Max	0.102297	0.817511	1.000000	-0.185224	
Rain	-0.044611	-0.049770	-0.185224	1.000000	
AM9_Temp	0.103578	0.900393	0.951034	-0.136218	
AM9_RH	-0.033659	-0.540743	-0.639255	0.220045	
AM9_Spd	-0.221265	-0.060050	-0.053351	0.076960	
PM3_Temp	0.022910	0.729600	0.907651	-0.160921	
PM3_RH	-0.052221	-0.520047	-0.808371	0.215631	
PM3_Spd	-0.232512	-0.063154	-0.022239	0.051807	
Total_Area	-0.045184	0.024080	-0.001396	-0.020690	
Bushfire_Event_Target	-0.040034	0.374604	0.330619	-0.084738	

	AM9_Temp	AM9_RH	AM9_Spd	PM3_Temp	PM3_RH	\
Station_Region	0.103578	-0.033659	-0.221265	0.022910	-0.052221	
Temps_Min	0.900393	-0.540743	-0.060050	0.729600	-0.520047	
Temps_Max	0.951034	-0.639255	-0.053351	0.907651	-0.808371	
Rain	-0.136218	0.220045	0.076960	-0.160921	0.215631	
AM9_Temp	1.000000	-0.693031	-0.060826	0.855122	-0.704679	
AM9_RH	-0.693031	1.000000	0.096861	-0.539372	0.656707	
AM9_Spd	-0.060826	0.096861	1.000000	0.317872	-0.206275	
PM3_Temp	0.855122	-0.539372	0.317872	1.000000	-0.852059	
PM3_RH	-0.704679	0.656707	-0.206275	-0.852059	1.000000	
PM3_Spd	-0.048645	0.101668	0.908667	0.316223	-0.208112	
Total_Area	0.029026	-0.165727	-0.005684	-0.011271	-0.027407	
Bushfire_Event_Target	0.359640	-0.404650	0.006320	0.286659	-0.329608	

	PM3_Spd	Total_Area	Bushfire_Event_Target
Station_Region	-0.232512	-0.045184	-0.040034
Temps_Min	-0.063154	0.024080	0.374604
Temps_Max	-0.022239	-0.001396	0.330619
Rain	0.051807	-0.020690	-0.084738
AM9_Temp	-0.048645	0.029026	0.359640
AM9_RH	0.101668	-0.165727	-0.404650
AM9_Spd	0.908667	-0.005684	0.006320
PM3_Temp	0.316223	-0.011271	0.286659
PM3_RH	-0.208112	-0.027407	-0.329608
PM3_Spd	1.000000	-0.007472	-0.010300
Total_Area	-0.007472	1.000000	0.117251
Bushfire_Event_Target	-0.010300	0.117251	1.000000



```
In [ ]: # Write DataFrame to CSV
csv_file_path = 'BB_R_Model_Data.csv' # Specify the path to the output CSV file
df.to_csv(csv_file_path, index=False) # Set index=False to avoid writing row indices
```

```
In [26]: features = ['Station_Region', 'Temps_Min', 'Temps_Max', 'Rain', 'AM9_Temp', 'AM9_RH', 'AM9_Spd', 'PM3_Temp', 'PM3_RH', 'PM3_Spd', 'Total_Area']
#features = ['Station_Region', 'Temps_Min', 'Temps_Max', 'Rain', 'AM9_RH', 'AM9_Spd']
X = df.loc[:, features]
y = df.loc[:, ['Bushfire_Event_Target']]
```

```
# Perform stratified split to preserve the class distribution
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
```

```
# Check the distribution of target variable in training and test sets
print("Train set class distribution:\n", y_train.value_counts())
print("Test set class distribution:\n", y_test.value_counts())
```

```
Train set class distribution:
Bushfire_Event_Target
```

```
0          160
1          112
```

```
Name: count, dtype: int64
```

```
Test set class distribution:
Bushfire_Event_Target
```

```
0           40
1           28
```

```
Name: count, dtype: int64
```

```
In [27]: X_train.head()
```

Out[27]:

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_Temp	AM9_RH	AM9_Spd	PM3_Temp
11223	6	14.5	18.5	0.4	15.9	62	28.0	15.0
2818	9	10.7	32.1	0.0	20.9	50	4.0	31.0
9194	4	18.0	22.6	0.0	19.3	37	24.0	22.0
13072	7	16.5	21.6	0.0	17.4	48	35.0	20.0
6862	8	8.4	27.5	0.0	16.6	58	9.0	26.0

In [28]:

X_test.head()

Out[28]:

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_Temp	AM9_RH	AM9_Spd	PM3_Temp
13265	7	3.8	17.8	0.0	11.0	99	15.0	17.0
16594	0	16.3	31.2	1.2	20.1	75	4.0	30.0
4151	10	9.8	28.3	0.0	17.3	70	9.0	27.0
7502	8	9.9	25.3	0.0	18.2	63	7.0	23.0
13081	7	10.8	34.7	0.0	25.7	31	15.0	33.0

In [29]:

y_train.head()

Out[29]:

	Bushfire_Event_Target
11223	1
2818	0
9194	0
13072	1
6862	1

In [30]:

X_test.info()

<class 'pandas.core.frame.DataFrame'>
Index: 68 entries, 13265 to 7492
Data columns (total 10 columns):
Column Non-Null Count Dtype
--- -
0 Station_Region 68 non-null int32
1 Temps_Min 68 non-null float64
2 Temps_Max 68 non-null float64
3 Rain 68 non-null float64
4 AM9_Temp 68 non-null float64
5 AM9_RH 68 non-null int64
6 AM9_Spd 68 non-null float64
7 PM3_Temp 68 non-null float64
8 PM3_RH 68 non-null int64
9 PM3_Spd 68 non-null float64
dtypes: float64(7), int32(1), int64(2)
memory usage: 5.6 KB

Random Forest Classifier

```
In [31]: # Initialize RandomForestClassifier with class_weight="balanced"
rf = RandomForestClassifier(class_weight='balanced', random_state=42)
```

```
In [32]: # Perform cross-validation
cv_scores = cross_val_score(rf, X_train, y_train.values.ravel(), cv=5)
print("Cross-validation scores:", cv_scores)
```

Cross-validation scores: [0.74545455 0.63636364 0.68518519 0.66666667 0.72222222]

```
In [33]: rf.fit(X_train, y_train)
```

C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\base.py:1151: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
return fit_method(estimator, *args, **kwargs)

```
Out[33]: RandomForestClassifier
RandomForestClassifier(class_weight='balanced', random_state=42)
```

```
In [34]: # Make predictions
y_pred = rf.predict(X_test)
```

```
In [35]: # Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", conf_matrix)

# Calculate precision, recall, and F1 score
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print precision, recall, and F1 score
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Accuracy: 0.6323529411764706

Confusion Matrix:

```
[[30 10]
 [15 13]]
```

Precision: 0.5652173913043478

Recall: 0.4642857142857143


F1 Score: 0.5098039215686274

```
In [36]: # Get feature importances
feature_importances = pd.DataFrame(rf.feature_importances_, index=X.columns, columns=['importance'])
feature_importances['abs_importance'] = feature_importances['importance'].abs()

# Filter feature importances based on absolute value
feature_filtered = feature_importances[feature_importances['abs_importance'] > 0.05]

# Plotting
plt.figure(figsize=(10, 6))
sns.barplot(x=feature_filtered['importance'], y=feature_filtered.index, palette='vibrant')
```

```
plt.title('Feature Importance from Random Forest Classifier')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.axvline(0, color='grey', linestyle='--') # Add a vertical line at x=0
plt.show()
```



Dimensionality Reduction

```
In [37]: features = ['Station_Region', 'Temps_Min', 'Temps_Max', 'Rain', 'AM9_RH', 'AM9_Spd', 'PM3_RH', 'PM3_Temp', 'AM9_Spd', 'PM3_Spd', 'Station_Region']
X = df.loc[:, features]
y = df.loc[:, ['Bushfire_Event_Target']]

# Perform stratified split to preserve the class distribution
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)

# Check the distribution of target variable in training and test sets
print("Train set class distribution:\n", y_train.value_counts())
print("Test set class distribution:\n", y_test.value_counts())
```

```
Train set class distribution:
Bushfire_Event_Target
0          160
1          112
Name: count, dtype: int64
Test set class distribution:
Bushfire_Event_Target
0           40
1           28
Name: count, dtype: int64
```

```
In [38]: # Initialize RandomForestClassifier with class_weight="balanced"
rf = RandomForestClassifier(class_weight='balanced', random_state=42)
```

```
In [39]: # Perform cross-validation
cv_scores = cross_val_score(rf, X_train, y_train.values.ravel(), cv=5)
print("Cross-validation scores:", cv_scores)
```

```
Cross-validation scores: [0.70909091 0.58181818 0.72222222 0.72222222 0.7962963 ]
```

In [40]: `rf.fit(X_train, y_train)`

C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\base.py:1151: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
 return fit_method(estimator, *args, **kwargs)

Out[40]:

▼ RandomForestClassifier
 RandomForestClassifier(class_weight='balanced', random_state=42)

In [41]: *# Evaluate model performance*
`accuracy = accuracy_score(y_test, y_pred)`
`print("Accuracy:", accuracy)`

Generate a confusion matrix
`conf_matrix = confusion_matrix(y_test, y_pred)`
`print("Confusion Matrix:\n", conf_matrix)`

Calculate precision, recall, and F1 score
`precision = precision_score(y_test, y_pred)`
`recall = recall_score(y_test, y_pred)`
`f1 = f1_score(y_test, y_pred)`

Print precision, recall, and F1 score
`print("Precision:", precision)`
`print("Recall:", recall)`
`print("F1 Score:", f1)`

Accuracy: 0.6323529411764706

Confusion Matrix:

```
[[30 10]
 [15 13]]
```

Precision: 0.5652173913043478

Recall: 0.4642857142857143

F1 Score: 0.5098039215686274

Dimesionality Reduced but the model performance is kept near the same level. By removing AM9_Temp majority of the model performance (F1 Score) increased as well.

Lositic regression with probability

In [784... *# Feature scaling (important for logistic regression)*
`scaler = StandardScaler()`
`X_train_scaled = scaler.fit_transform(X_train)`
`X_test_scaled = scaler.transform(X_test)`

In [785... *# Create a logistic regression model with class weights*
`log_reg = LogisticRegression(random_state=0, class_weight='balanced')`

Perform cross-validation
`cv_scores = cross_val_score(log_reg, X_train, y_train.values.ravel(), cv=5)`
`print("Cross-validation scores:", cv_scores)`

Train the model
`log_reg.fit(X_train_scaled, y_train)`

Cross-validation scores: [0.63636364 0.61818182 0.75925926 0.62962963 0.66666667]

```
C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py:46
0: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py:46
```

```
0: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1184: Da
```

```
taConversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

Out[785]:

```
LogisticRegression
LogisticRegression(class_weight='balanced', random_state=0)
```

In [786...

```
# Make predictions
y_pred = log_reg.predict(X_test_scaled)
```

In [787...

```
# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", conf_matrix)

# Calculate precision, recall, and F1 score
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print precision, recall, and F1 score
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

```
Accuracy: 0.6470588235294118
```

```
Confusion Matrix:
```

```
[[27 13]
```

```
 [11 17]]
```

```
Precision: 0.5666666666666667
```

```
Recall: 0.6071428571428571
```

```
F1 Score: 0.5862068965517241
```

In [788...

```
# Make predictions with probabilities
y_pred_proba = log_reg.predict_proba(X_test_scaled)

# Create a DataFrame to display test data with predicted probabilities
test_data_with_proba = pd.DataFrame(X_test, columns=features)
test_data_with_proba['Predicted_Class'] = log_reg.predict(X_test_scaled)
test_data_with_proba['Probability_Class_0'] = y_pred_proba[:, 0]
```

```
test_data_with_proba['Probability_Class_1'] = y_pred_proba[:, 1]

# Display the first few rows of the test data with probabilities
print(test_data_with_proba.head())
```

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_RH	AM9_Spd	PM3_RH	\
13265	7	3.8	17.8	0.0	99	15.0	65	
16594	0	16.3	31.2	1.2	75	4.0	27	
4151	10	9.8	28.3	0.0	70	9.0	26	
7502	8	9.9	25.3	0.0	63	7.0	39	
13081	7	10.8	34.7	0.0	31	15.0	16	

	Predicted_Class	Probability_Class_0	Probability_Class_1
13265	0	0.882791	0.117209
16594	1	0.312355	0.687645
4151	0	0.600457	0.399543
7502	0	0.515639	0.484361
13081	1	0.290802	0.709198

Decision Tree Classifier

```
In [789... tree_clf = DecisionTreeClassifier(random_state=0, class_weight='balanced')

# Perform cross-validation
cv_scores = cross_val_score(tree_clf, X_train, y_train.values.ravel(), cv=5)
print("Cross-validation scores:", cv_scores)

tree_clf.fit(X_train, y_train)

Cross-validation scores: [0.6          0.63636364 0.66666667 0.66666667 0.68518519]
```

```
Out[789]: DecisionTreeClassifier
DecisionTreeClassifier(class_weight='balanced', random_state=0)
```

```
In [790... # Make predictions
y_pred = tree_clf.predict(X_test)
```

```
In [791... # Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", conf_matrix)

# Calculate precision, recall, and F1 score
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print precision, recall, and F1 score
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```


Accuracy: 0.5735294117647058
 Confusion Matrix:
 [[30 10]
 [19 9]]
 Precision: 0.47368421052631576
 Recall: 0.32142857142857145
 F1 Score: 0.3829787234042553

Gradient Boosting Classifier

In [792]...

```
from sklearn.ensemble import GradientBoostingClassifier

gb_clf = GradientBoostingClassifier(random_state=0)

# Perform cross-validation
cv_scores = cross_val_score(gb_clf, X_train, y_train.values.ravel(), cv=5)
print("Cross-validation scores:", cv_scores)

gb_clf.fit(X_train, y_train)
```

Cross-validation scores: [0.65454545 0.61818182 0.68518519 0.72222222 0.75925926]

C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\ensemble_gb.py:424: DataCon
 versionWarning: A column-vector y was passed when a 1d array was expected. Please
 change the shape of y to (n_samples,), for example using ravel().
 y = column_or_1d(y, warn=True)

Out[792]:

▼ GradientBoostingClassifier
 GradientBoostingClassifier(random_state=0)

In []:

```
# Make predictions
y_pred = gb_clf.predict(X_test)
```

In [793]...

```
# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", conf_matrix)

# Calculate precision, recall, and F1 score
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print precision, recall, and F1 score
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Accuracy: 0.5735294117647058
 Confusion Matrix:
 [[30 10]
 [19 9]]
 Precision: 0.47368421052631576
 Recall: 0.32142857142857145
 F1 Score: 0.3829787234042553

Support Vector Machine (SVM)

```
In [42]: svc_clf = SVC(random_state=0, class_weight='balanced', probability=True)

# Perform cross-validation
cv_scores = cross_val_score(svc_clf, X_train, y_train.values.ravel(), cv=5)
print("Cross-validation scores:", cv_scores)

svc_clf.fit(X_train, y_train)
```

Cross-validation scores: [0.69090909 0.58181818 0.7962963 0.62962963 0.62962963]

C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

```
Out[42]: SVC
SVC(class_weight='balanced', probability=True, random_state=0)
```

```
In [43]: X_train.head(10)
```

```
Out[43]:
```

	Station_Region	Temps_Min	Temps_Max	Rain	AM9_RH	AM9_Spd	PM3_RH
11223	6	14.5	18.5	0.4	62	28.0	83
2818	9	10.7	32.1	0.0	50	4.0	19
9194	4	18.0	22.6	0.0	37	24.0	37
13072	7	16.5	21.6	0.0	48	35.0	20
6862	8	8.4	27.5	0.0	58	9.0	26
3506	10	14.9	31.8	0.0	66	11.0	34
18317	5	18.0	30.6	0.0	48	7.0	29
17250	0	7.8	23.7	0.0	80	1032.1	2
10468	6	-1.3	7.3	14.6	100	20.0	83
19275	5	9.2	20.7	0.2	61	6.0	61

```
In [44]: # Make predictions
y_pred = svc_clf.predict(X_test)
```

```
In [45]: # Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", conf_matrix)

# Calculate precision, recall, and F1 score
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print precision, recall, and F1 score
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Accuracy: 0.6764705882352942
 Confusion Matrix:
 [[23 17]
 [5 23]]
 Precision: 0.575
 Recall: 0.8214285714285714
 F1 Score: 0.676470588235294

In [46]:

```
# Get model parameters
model_params = svc_clf.get_params()
print("Model parameters:", model_params)
```

Model parameters: {'C': 1.0, 'break_ties': False, 'cache_size': 200, 'class_weight': 'balanced', 'coef0': 0.0, 'decision_function_shape': 'ovr', 'degree': 3, 'gamma': 'scale', 'kernel': 'rbf', 'max_iter': -1, 'probability': True, 'random_state': 0, 'shrinking': True, 'tol': 0.001, 'verbose': False}

k-Nearest Neighbors (k-NN)

In [797]:

```
knn_clf = KNeighborsClassifier()

# Perform cross-validation
cv_scores = cross_val_score(knn_clf, X_train, y_train.values.ravel(), cv=5)
print("Cross-validation scores:", cv_scores)

knn_clf.fit(X_train, y_train)
```

Cross-validation scores: [0.63636364 0.6 0.68518519 0.68518519 0.72222222]

C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\neighbors_classification.py:228: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return self._fit(X, y)
```

Out[797]:

```
▼ KNeighborsClassifier
KNeighborsClassifier()
```

In [798]:

```
# Make predictions
y_pred = knn_clf.predict(X_test)
```

In [799]:

```
# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", conf_matrix)

# Calculate precision, recall, and F1 score
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print precision, recall, and F1 score
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Accuracy: 0.6029411764705882
 Confusion Matrix:
 [[30 10]
 [17 11]]
 Precision: 0.5238095238095238
 Recall: 0.39285714285714285
 F1 Score: 0.4489795918367347

Naive Bayes

```
In [800]: nb_clf = GaussianNB()

# Perform cross-validation
cv_scores = cross_val_score(nb_clf, X_train, y_train.values.ravel(), cv=5)
print("Cross-validation scores:", cv_scores)

nb_clf.fit(X_train, y_train)
```

Cross-validation scores: [0.63636364 0.61818182 0.75925926 0.59259259 0.57407407]
 C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
 y = column_or_1d(y, warn=True)

```
Out[800]: ▾ GaussianNB
GaussianNB()
```

```
In [801]: # Make predictions
y_pred = nb_clf.predict(X_test)
```

```
In [802]: # Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", conf_matrix)

# Calculate precision, recall, and F1 score
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print precision, recall, and F1 score
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Accuracy: 0.6470588235294118
 Confusion Matrix:
 [[22 18]
 [6 22]]
 Precision: 0.55
 Recall: 0.7857142857142857
 F1 Score: 0.6470588235294117

Neural Network

In [803...

!

In [804...

```
# Import necessary libraries
import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from sklearn.metrics import accuracy_score
```

In [805...

```
# Reshape input for RNN (samples, time_steps, features)
# Assuming time_steps = 1 for non-sequential data, you can adjust if your data is s
X_train_rnn = X_train_scaled.reshape(X_train_scaled.shape[0], 1, X_train_scaled.sha
X_test_rnn = X_test_scaled.reshape(X_test_scaled.shape[0], 1, X_test_scaled.shape[1
```

In [806...

```
# Build the RNN model
model = Sequential()
```

In [807...

```
# Add LSTM Layer
model.add(LSTM(50, activation='relu', input_shape=(X_train_rnn.shape[1], X_train_rnn
```

C:\Users\Thinithi\anaconda3\Lib\site-packages\keras\src\layers\rnn\rnn.py:204: Use
rWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using
Sequential models, prefer using an `Input(shape)` object as the first layer in the
model instead.
super().__init__(**kwargs)

In [808...

```
# Add a dropout layer to prevent overfitting
model.add(Dropout(0.2))
```

In [809...


```
# Output layer for binary classification
model.add(Dense(1, activation='sigmoid'))
```


In [810...


```
# Compile the model with optimizer, loss function, and evaluation metric
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```


In [811...


```
# Train the model with training data
model.fit(X_train_rnn, y_train, epochs=50, batch_size=32, validation_split=0.1, ver
```


Epoch 1/50
8/8  2s 35ms/step - accuracy: 0.5214 - loss: 0.6869 - val_accuracy: 0.6429 - val_loss: 0.6811


Epoch 2/50
8/8  0s 6ms/step - accuracy: 0.6366 - loss: 0.6773 - val_accuracy: 0.7500 - val_loss: 0.6724


Epoch 3/50
8/8  0s 8ms/step - accuracy: 0.6752 - loss: 0.6676 - val_accuracy: 0.6786 - val_loss: 0.6638


Epoch 4/50
8/8  0s 7ms/step - accuracy: 0.6333 - loss: 0.6686 - val_accuracy: 0.6786 - val_loss: 0.6557


Epoch 5/50
8/8  0s 6ms/step - accuracy: 0.6919 - loss: 0.6541 - val_accuracy: 0.6786 - val_loss: 0.6478


Epoch 6/50
8/8  0s 7ms/step - accuracy: 0.7223 - loss: 0.6364 - val_accuracy: 0.6786 - val_loss: 0.6399


Epoch 7/50
8/8  0s 6ms/step - accuracy: 0.6617 - loss: 0.6414 - val_accuracy: 0.6786 - val_loss: 0.6323


Epoch 8/50
8/8  0s 6ms/step - accuracy: 0.6756 - loss: 0.6261 - val_accuracy: 0.6786 - val_loss: 0.6238


Epoch 9/50
8/8  0s 6ms/step - accuracy: 0.6937 - loss: 0.6212 - val_accuracy: 0.6786 - val_loss: 0.6148

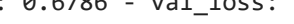
Epoch 10/50
8/8  0s 9ms/step - accuracy: 0.6509 - loss: 0.6236 - val_accuracy: 0.6786 - val_loss: 0.6062

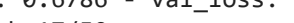
Epoch 11/50
8/8  0s 6ms/step - accuracy: 0.6883 - loss: 0.6110 - val_accuracy: 0.6786 - val_loss: 0.5987


Epoch 12/50
8/8  0s 6ms/step - accuracy: 0.6805 - loss: 0.6028 - val_accuracy: 0.6786 - val_loss: 0.5906


Epoch 13/50
8/8  0s 6ms/step - accuracy: 0.7162 - loss: 0.5747 - val_accuracy: 0.6786 - val_loss: 0.5837


Epoch 14/50
8/8  0s 6ms/step - accuracy: 0.7225 - loss: 0.5721 - val_accuracy: 0.6786 - val_loss: 0.5778


Epoch 15/50
8/8  0s 5ms/step - accuracy: 0.7251 - loss: 0.5708 - val_accuracy: 0.6786 - val_loss: 0.5714


Epoch 16/50
8/8  0s 6ms/step - accuracy: 0.7190 - loss: 0.5558 - val_accuracy: 0.6786 - val_loss: 0.5671

Epoch 17/50
8/8  0s 6ms/step - accuracy: 0.6944 - loss: 0.5686 - val_accuracy: 0.6786 - val_loss: 0.5624


Epoch 18/50
8/8  0s 5ms/step - accuracy: 0.7176 - loss: 0.5602 - val_accuracy: 0.6786 - val_loss: 0.5586


Epoch 19/50
8/8  0s 6ms/step - accuracy: 0.6664 - loss: 0.5800 - val_accuracy: 0.6786 - val_loss: 0.5544


Epoch 20/50
8/8  0s 6ms/step - accuracy: 0.7134 - loss: 0.5382 - val_accuracy: 0.6786 - val_loss: 0.5511


Epoch 21/50
8/8  0s 6ms/step - accuracy: 0.6496 - loss: 0.5748 - val_accuracy: 0.6786 - val_loss: 0.5494


Epoch 22/50


8/8  0s 5ms/step - accuracy: 0.7002 - loss: 0.5290 - val_accu
acy: 0.6786 - val_loss: 0.5469
Epoch 23/50


8/8  0s 7ms/step - accuracy: 0.6857 - loss: 0.5493 - val_accu
acy: 0.6429 - val_loss: 0.5446
Epoch 24/50


8/8  0s 8ms/step - accuracy: 0.7141 - loss: 0.5369 - val_accu
acy: 0.6429 - val_loss: 0.5442
Epoch 25/50


8/8  0s 6ms/step - accuracy: 0.7192 - loss: 0.5369 - val_accu
acy: 0.6429 - val_loss: 0.5430
Epoch 26/50


8/8  0s 6ms/step - accuracy: 0.6812 - loss: 0.5782 - val_accu
acy: 0.6786 - val_loss: 0.5421
Epoch 27/50


8/8  0s 6ms/step - accuracy: 0.6729 - loss: 0.5773 - val_accu
acy: 0.6786 - val_loss: 0.5412
Epoch 28/50


8/8  0s 6ms/step - accuracy: 0.7041 - loss: 0.5471 - val_accu
acy: 0.6786 - val_loss: 0.5408
Epoch 29/50

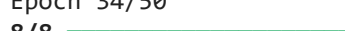
8/8  0s 6ms/step - accuracy: 0.6696 - loss: 0.5467 - val_accu
acy: 0.6786 - val_loss: 0.5390
Epoch 30/50


8/8  0s 6ms/step - accuracy: 0.7234 - loss: 0.5178 - val_accu
acy: 0.6786 - val_loss: 0.5387
Epoch 31/50


8/8  0s 6ms/step - accuracy: 0.6994 - loss: 0.5397 - val_accu
acy: 0.6786 - val_loss: 0.5373
Epoch 32/50


8/8  0s 6ms/step - accuracy: 0.6958 - loss: 0.5714 - val_accu
acy: 0.6786 - val_loss: 0.5366
Epoch 33/50


8/8  0s 5ms/step - accuracy: 0.6801 - loss: 0.5347 - val_accu
acy: 0.6786 - val_loss: 0.5356
Epoch 34/50


8/8  0s 7ms/step - accuracy: 0.6789 - loss: 0.5449 - val_accu
acy: 0.6429 - val_loss: 0.5336
Epoch 35/50


8/8  0s 5ms/step - accuracy: 0.7073 - loss: 0.5307 - val_accu
acy: 0.6429 - val_loss: 0.5325
Epoch 36/50


8/8  0s 6ms/step - accuracy: 0.6922 - loss: 0.5222 - val_accu
acy: 0.6429 - val_loss: 0.5320
Epoch 37/50


8/8  0s 8ms/step - accuracy: 0.7108 - loss: 0.5247 - val_accu
acy: 0.6429 - val_loss: 0.5317
Epoch 38/50


8/8  0s 5ms/step - accuracy: 0.6993 - loss: 0.5306 - val_accu
acy: 0.6429 - val_loss: 0.5317
Epoch 39/50

8/8  0s 5ms/step - accuracy: 0.6940 - loss: 0.5204 - val_accu
acy: 0.6429 - val_loss: 0.5306
Epoch 40/50

8/8  0s 6ms/step - accuracy: 0.6607 - loss: 0.5692 - val_accu
acy: 0.6071 - val_loss: 0.5289
Epoch 41/50

8/8  0s 6ms/step - accuracy: 0.7115 - loss: 0.5350 - val_accu
acy: 0.6071 - val_loss: 0.5276
Epoch 42/50

8/8  0s 8ms/step - accuracy: 0.6637 - loss: 0.5549 - val_accu
acy: 0.6071 - val_loss: 0.5258
Epoch 43/50

8/8  0s 5ms/step - accuracy: 0.7129 - loss: 0.5223 - val_accu

acy: 0.6071 - val_loss: 0.5260

Epoch 44/50

8/8 ————— 0s 6ms/step - accuracy: 0.7380 - loss: 0.5169 - val_accu

acy: 0.6071 - val_loss: 0.5234

Epoch 45/50

8/8 ————— 0s 6ms/step - accuracy: 0.6716 - loss: 0.5583 - val_accu

acy: 0.6071 - val_loss: 0.5222

Epoch 46/50

8/8 ————— 0s 5ms/step - accuracy: 0.6276 - loss: 0.5647 - val_accu

acy: 0.6071 - val_loss: 0.5207

Epoch 47/50

8/8 ————— 0s 5ms/step - accuracy: 0.7199 - loss: 0.4986 - val_accu

acy: 0.6071 - val_loss: 0.5190

Epoch 48/50

8/8 ————— 0s 6ms/step - accuracy: 0.7064 - loss: 0.5189 - val_accu

acy: 0.6429 - val_loss: 0.5165

Epoch 49/50

8/8 ————— 0s 6ms/step - accuracy: 0.6878 - loss: 0.5158 - val_accu

acy: 0.6429 - val_loss: 0.5167

Epoch 50/50

8/8 ————— 0s 5ms/step - accuracy: 0.6812 - loss: 0.5474 - val_accu

acy: 0.6429 - val_loss: 0.5166

<keras.src.callbacks.history.History at 0x2ba51dece90>

Out[811]:

In [812...]

```
# Make predictions on the test set
y_pred_prob = model.predict(X_test_rnn)
y_pred = (y_pred_prob > 0.5).astype("int32") # Convert probabilities to 0 or 1
```

3/3 ————— 0s 59ms/step

In [813...]

```
# Calculate accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

# Calculate precision, recall, and F1 score
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print precision, recall, and F1 score
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Accuracy: 0.6176470588235294

Precision: 0.5454545454545454

Recall: 0.42857142857142855

F1 Score: 0.47999999999999999

Model Evaluation

In [817...]

```
# Load the image
img = Image.open("Model Evaluation.png") # replace with your image path

# Create a larger figure
plt.figure(figsize=(10, 10))

# Display the image
plt.imshow(img)
plt.axis('off') # Turn off axis labels
plt.show()
```


Classifier	Accuracy	Confusion Matrix	Precision	Recall	F1 Score
Random Forest Classifier	0.6324	[[30, 10], [15, 13]]	0.5652	0.4643	0.5098
Logistic Regression	0.6471	[[27, 13], [11, 17]]	0.5667	0.6071	0.5862
Decision Tree Classifier	0.5735	[[30, 10], [19, 9]]	0.4737	0.3214	0.383
Gaussian Boosting Classifier	0.5735	[[30, 10], [19, 9]]	0.4737	0.3214	0.383
Support Vector Machine	0.6765	[[23, 17], [5, 23]]	0.575	0.8214	0.6765
KNN Classifier	0.6029	[[30, 10], [17, 11]]	0.5238	0.3929	0.449
Naive Bayes Classifier	0.6471	[[22, 18], [6, 22]]	0.55	0.7857	0.6471
RNN Classifier	0.6176	Not provided	0.5455	0.4286	0.48

```
In [ ]: # Save the model
joblib.dump(svc_clf, 'BB_svm_model.pkl')
```

```
In [6]: # Load the model
loaded_model = joblib.load('BB_svm_model.pkl')
```

```
In [7]: X_test = [9,10.7,32.1,0.0,23,4.0,15]

# Reshape it to a 2D array
X_test = np.array(X_test).reshape(1, -1)

# Use the Loaded model to make predictions
y_pred = loaded_model.predict(X_test)
print(y_pred)
```

```
[1]
```

C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but SVC was fitted with feature names
warnings.warn(

```
In [11]: # Get probability estimates
probabilities = loaded_model.predict_proba(X_test)
print("Probabilities:", probabilities)
print(probabilities[:,1])
```

```
Probabilities: [[0.29005412 0.70994588]]
[0.70994588]
```

C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but SVC was fitted with feature names
warnings.warn(

```
In [12]: X_test = [5,9.2,20.7,0.2,61,6.0,61]

# Reshape it to a 2D array
X_test = np.array(X_test).reshape(1, -1)

# Use the Loaded model to make predictions
y_pred = loaded_model.predict(X_test)
print(y_pred)
```

```
[0]
```

C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but SVC was fitted with feature names
warnings.warn(

```
In [14]: # Get probability estimates
probabilities = loaded_model.predict_proba(X_test)
print("Probabilities:", probabilities)
print(probabilities[:,0])
```

```
Probabilities: [[0.63029646 0.36970354]]
[0.63029646]
```

```
C:\Users\Thinithi\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
warnings.warn(
```

Hyper Parameter Tuning

```
In [43]: # Define the parameter grid for tuning
param_grid = {
    'C': [0.1, 1, 10, 100], # Regularization parameter
    'gamma': ['scale', 'auto', 0.001, 0.01, 0.1], # Kernel coefficient
    'kernel': ['rbf', 'linear'], # Types of kernel
    'degree': [2, 3, 4] # Degree for polynomial kernel
}
```

```
In [44]: # Initialize the SVC classifier
svc_clf = SVC(class_weight='balanced', probability=True, random_state=0)
```

```
In [ ]: # Initialize GridSearchCV
grid_search = GridSearchCV(estimator=svc_clf, param_grid=param_grid,
                           scoring='accuracy', cv=5, n_jobs=-1, verbose=1)

# Fit GridSearchCV
grid_search.fit(X_train, y_train.values.ravel())

# Retrieve the best parameters
best_params = grid_search.best_params_
best_score = grid_search.best_score_

print("Best Parameters:", best_params)
print("Best Cross-Validation Score:", best_score)
```

Fitting 5 folds for each of 120 candidates, totalling 600 fits

```
In [ ]:
```